

# Cloud Computing Labs: Using AWS and Docker

The primary purpose of this assignment is to get familiar with the most popular Infrastructure-as-a-Service: Amazon Web Services and the container technique: Docker. This assignment requires a lot of activities on the command line and some coding work with Python.

## Part 1: Setup AWS and Docker

1.1 Review the lecture about using AWS. Create your own AWS account. Note that a credit card will be needed for creating the account. New users will get the [AWS Free Tier](#) for one year, which will be sufficient for this lab. Or you can also use the department's [Amazon Educate resource](#) to get the free credits.

- Create your AWS access keys. Go to the menu "My Security Credentials" and find "Access keys" to create the Access Key ID and Secret Access Key.
- Setup AWS command line tool. Use the following commands in your ubuntu box.

```
> sudo apt-get install python-pip
> sudo pip install awscli
```

use "aws help" to explore the command usage. Config the aws tool

```
> aws configure
AWS Access Key ID [None]: XXXXXXXX
AWS Secret Access Key [None]: XXXXXXXX
Default region name [None]: us-east-1
Default output format [None]: json
```

- Create a security group.

```
> aws ec2 create-security-group --group-name your_security_group_name --descripti
> aws ec2 authorize-security-group-ingress --group-name your_security_group_name
```

- Create a key pair.

```
> aws ec2 create-key-pair --key-name your_key_name
--query 'KeyMaterial' --output text > your_key_name.pem
> chmod 400 your_key_name.pem
```

- Some instances may need to specify the subnet id. Please check [VPC and subnets](#) or use "aws ec2 describe-subnets" to find the subnet information.
- Read the slides and the [Python boto3 library documentation](#) and understand how to use boto3 to access EC2 and S3.

1.2 setup Docker in your linux system.

Please check the lecture slides and the [online tutorial](#) of Docker and learn how to setup the Docker in linux systems and use the dockerfile method for creating Docker images.

Now, answer the following questions:

**Question 1.1** create an instance with the commands discussed in the lecture note. Use free-tier images (e.g., ami-cd0f5cb6) if you are using the "free tier" resource. Remember that for free tier instances, you also need to use "--instance-type t2.micro" and specify the subnet id with "--subnet-id". Finally login the instance with ssh. Save your screen shots to show that you have successfully done.

**Question 1.2** Use the boto APIs to implement a python function `start_instances(num_instances)`, where the parameter `num_instances` is the number of instances you will be creating. This function will create a number of instances and wait until the state of the instances become "running", and then return the list of instance ids.

**Question 1.3** Write a python script that uses the boto APIs to find out all the files in the bucket "wsu2017fall", print out the contents in the files, and copy the files to your own bucket. Remember to handle exceptions (e.g., empty directories). Keep your bucket undeleted until we finish grading!

**Question 1.4** Create a Docker image based on ubuntu image. Let's assume a scenario that you can remotely login a running container and debug a pyspark script. Therefore, the image should contain a ssh server, the single-node Spark setup, and a simple pyspark script (e.g., the wordcount program). (1) Post your dockerfile and your Docker hub link. (2) Post the command that starts the container and exposes its ssh port to external via the host's 2222 port (hint: check the `-p` option). (3) Post the screenshot showing that you can remotely login the container in an AWS instance and test-run it, e.g., "spark-submit wordcount.py" successfully.

**Question 1.5** How much time did you spend on these tasks? (hours)

**Question 1.6** How useful is this task to your understanding of EC2 and S3? (very useful, somewhat useful, not useful)

## Part 2: Monitoring VM instances and Docker containers

**Question 2.1** In this task, you will implement a tool with Python Boto3 library and the [Paramiko](#) Python SSH library to monitor the status of the instances you created. This monitoring tool will constantly (e.g., every 5 seconds) print out the CPU usage of each instance. Note that you can execute commands in instances remotely via ssh, like

```
ssh -i your_private_key.pem ubuntu@EC2_instance_Public_DNS "top -bn1 | grep Cpu"
```

The command "top -bn1 | grep Cpu" will get the the line of the command "top" output that contains Cpu information. The output of the remote command execution will be sent to you.

In your python code, you will need to create 2 instances using the function created in Q1.2 and then in a loop every 5 seconds the command is executed remotely in the instances by using the ssh functions provided by the Paramiko library, and print out the information "instance\_ id \t Cpu usage".

**Question 2.2** Extend your tool to monitor Docker containers in VM instances. Assume you have started 2 EC2 instances using the python function you developed. It's better to use an image that contains Docker. If not, for each EC2 instance, you can install Docker manually or via ssh commands in your python script. Then, use ssh command in python to start 2 Docker container daemons as follows (e.g., using the ubuntu image). Note that the `-d` option is used to run the container as a daemon.

```
docker run -d -t ubuntu sh
```

You can retrieve the container IDs (similar to VM instance IDs) using the following or other similar command.

```
docker ps | grep ubuntu
```

To execute a command in the container, for instance, getting the CPU usage, you can use

```
docker exec container_ID top -bn1 | grep CPU
```

Now you implement your python program to monitor the CPU usage of each container in each instance every 5 seconds and print out "instance\_ID \t container\_ID \t CPU usage".

After you implement these tools, please answer the following questions:

**Question 2.3** How much time did you spend on this part? (hours)

**Question 2.4** How useful is this task to your understanding of controlling AWS and Docker in programs? (very useful, somewhat useful, not useful)

## Final Survey Questions

**Question 3.1** Your level of interest in this lab exercise (high, average, low);

**Question 3.2** How challenging is this lab exercise? (high, average, low);

**Question 3.3** How valuable is this lab as a part of the course (high, average, low);

**Question 3.4** Are the supporting materials and lectures helpful for you to finish the project? (very helpful, somewhat helpful, not helpful);

**Question 3.5** Do you feel confident on applying the skills learned in the lab to solve other problems with AWS EC2 and S3?

## Deliverables

Turn in (1) the code and answers for Questions 1.1-1.4 and 2.1-2.2 in one PDF file. No need to rephrase the questions. For coding problems, you should copy&paste your code in text. Do not post screenshots of your code or answers, unless you are asked to do so. (2) All survey questions go to the second PDF. DO NOT zip the two files. Without following the instructions, you will get 0 for those questions immediately.

**Make sure that you have terminated all instances after finishing your work!**

---

This page, first created: Nov. 1, 2016; last updated: Oct. 2019